



BEYOND CLASSICAL SEARCH

Fundamentals of Artificial Intelligence

Session 09

Pramod Sharma

pramod.sharma@prasami.com

2

Agenda

- Local Search Algorithms
- Hill-climbing
- Simulated Annealing
- Local Beam Search
- Genetic Algorithms
- Parallel Search

4/11/2024

pra-sâmi

3

Search Algorithms So Far

- ❑ Designed to explore search space systematically
- ❑ Keep one or more paths in memory
- ❑ Record which have been explored and which have not
- ❑ A path to goal represents the solution
- ❑ More often than not, are complex in terms on time and space

4/11/2024

pra-sâmi

4

Local Search Algorithms

- ❑ In many optimization problems, the path to the goal is irrelevant; the goal state itself is the solution
- ❑ State space = set of "complete" configurations
- ❑ Find configuration satisfying constraints,
 - ❖ e.g., n-queens, factory floor layout, job shop scheduling, vehicle routing and portfolio management
- ❑ In such cases, we can use **local search algorithms**
- ❑ Keep a single "current" state, try to improve it
 - ❖ Use very little memory – usually a constant amount
 - ❖ Find reasonable solutions in large or infinite state spaces for which systematic solutions are unsuitable
 - ❖ Useful for solving optimization problems, e.g. Darwinian evolution, no “goal test” or “path cost”

4/11/2024

pra-sâmi


5

Example: n-Queen Problem

- ❑ Put n queens on an $n \times n$ board with no two queens under attack
 - ❖ Not on the same row, column, or diagonal



- ❑ We still have one conflict
 - ❖ This is best we could do in present search



- ❑ Good neighborhood function is good balance between immediate neighbors and length of path to solution.
 - ❖ It needs to be learned. Cannot be guessed!

4/11/2024

pra-sâmi

6

Search Space

- ❑ State
 - ❖ All queen on board in some configuration
- ❑ Successor function
 - ❖ Move single queen to another square in the same column
- ❑ Objective function
 - ❖ No of queens attacking each other
- ❑ We are keen in finding a state which minimizes is objective function

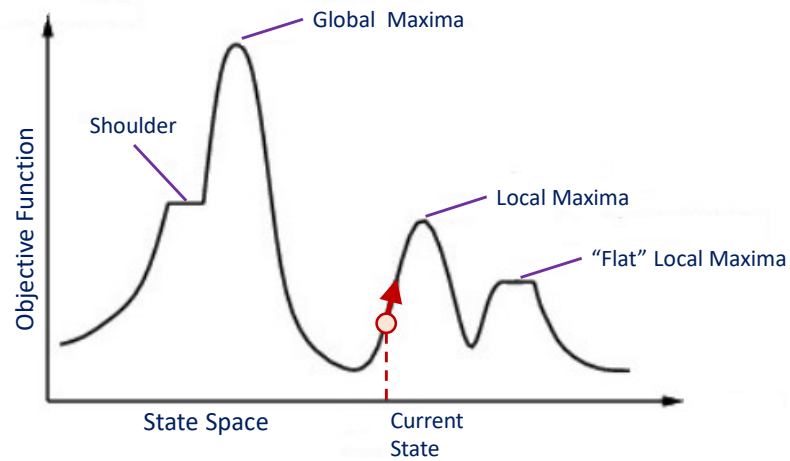
4/11/2024

pra-sâmi

7

State Space Landscape

- Problem: depending on initial state, can get stuck in local maxima/minima



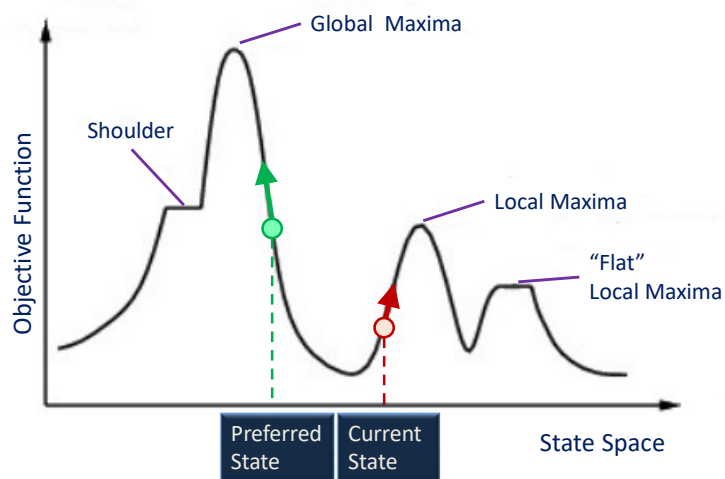
4/11/2024

pra-sâmi


8

State Space Landscape

- Problem: depending on initial state, can get stuck in local maxima/minima



4/11/2024


Local Minimum with $h = 1$

pra-sâmi

9

Example: Initial State

- Assume the objective function measures the straight-line distance




4/11/2024

pra-sâmi

10

Example: Local Minima

- Assume the objective function measures the straight-line distance




4/11/2024

pra-sâmi

11

Example: A Plausible Solution

- Assume the objective function measure the straight-line distance
- Making some “bad” choices is actually not that bad




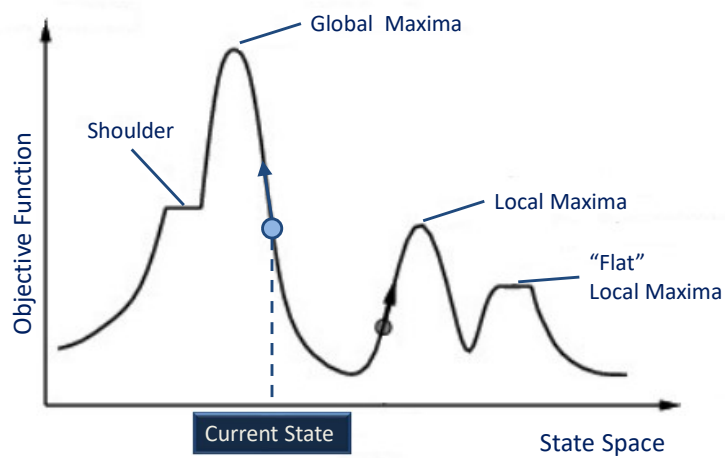
4/11/2024

pra-sâmi

12

State Space Landscape

- Right choice will take you to global maxima



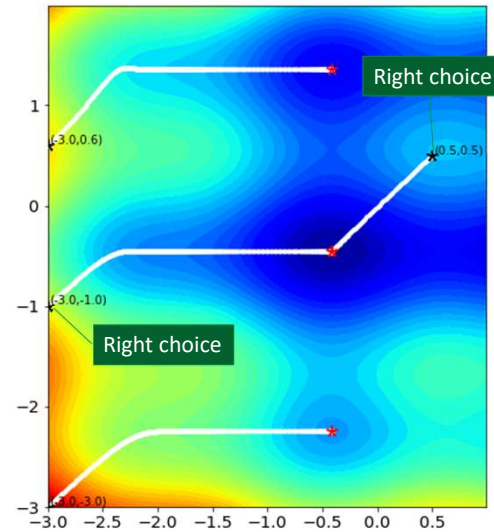


Global Maximum

4/11/2024

pra-sâmi

13

State Space Landscape






4/11/2024

pra-sâmi

14

State Space Landscape



4/11/2024

pra-sâmi

15

Hill Climbing



4/11/2024

pra-sâmi


16

Iterative Improvement

- Consider all states laid out on the surface of a landscape



$$f(x,y) = e^{-(x^2+y^2)} + 2e^{-((x-1.7)^2 + (y-1.7)^2)}$$



- The height at any point corresponds to the result of the evaluation function

4/11/2024

pra-sâmi

17

Iterative Improvement

- ❑ Paths typically not retained - very little memory needed
- ❑ Move around the landscape trying to find the highest peaks – optimal solutions (or lowest valleys the if trying to minimise)
 - ❖ Useful for hard, practical problems where the state description itself holds all the information needed for a solution
 - ❖ Find reasonable solutions in a large or infinite state space
- ❑ Example : travelling salesperson, neural network gradient descent,
 - ❖ After back propagation, nudge it a bit and see if it converges better

4/11/2024

pra-sâmi

18

Iterative Improvement

- ❑ Consider all states laid out on the surface of a landscape


- ❑ Random Sampling

- ❖ Generate a state randomly and compare

- ❑ Random walk

- ❖ Randomly pick neighbour of current state

- ❑ The height at any point corresponds to the result of the evaluation function



$$f(x,y) = e^{-(x^2+y^2)} + 2e^{-((x-1.7)^2 + (y-1.7)^2)}$$




4/11/2024

pra-sâmi

19

Hill-Climbing Search

- ❑ Check all neighbours find better and move towards that neighbour
 - ❖ Terminate when peak is reached
- ❑ Maximize objective function
- ❑ Never thinks beyond its neighbours
- ❑ Can randomly choose among the best successors
 - ❖ If multiple successors have best value




4/11/2024

pra-sâmi

20

Hill-Climbing Search

- ❑ “Like climbing Everest in thick fog with amnesia”
- ❑ Complete? Optimal?
- ❑ Hill climbing search is sometimes called **greedy local search**
- ❑ Although greedy algorithms often perform well, hill climbing gets stuck when:
 - ❖ Local maxima/minima
 - ❖ Ridges
 - ❖ Plateau (shoulder or flat local maxima/minima)
- ❑ The steepest-ascent hill climbing solves only 14% of the randomly-generated 8-queen problems with an avg. of 4 steps
 - ❖ When it gets stuck (86% generated problems), it takes only 3 moves
- ❑ Allowing sideways move raises the success rate to 94% with an avg. of 21 steps, and 64 steps for each failure



4/11/2024

pra-sâmi

21

Hill- Climbing (Greedy Local)

- ❑ Start with current-state = initial-state
- ❑ Until current-state = goal-state OR there is no change in current-state do:
 - ❖ Get the children of current-state and apply evaluation function to each child
 - ❖ If one of the children has a better score, then set current-state to the child with the best score
- ❑ Loop that moves in the direction of increasing (or decreasing) value
 - ❖ Terminates when a “peak” (or “dip”) is reached
 - ❖ If more than one best direction, the algorithm can choose at random


4/11/2024

pra-sâmi


22

Hill Climbing: n-Queen Problem

- ❑ Move one queen at a time to reduce h



Eight Queen with heuristic cost estimate $h = 17$



Local Minimum with $h = 1$

4/11/2024

Only current state matters. How we got there is of little consequence.

pra-sâmi

23

"A hill climbing algorithm that never makes "downhill" (or "uphill") to a lower (or "higher") value, Does not guarantee complete search!

A purely random walk – moving to a successor chosen uniformly at random – is complete, but extremely inefficient!



"

What should we do?


4/11/2024

pra-sâmi

24

Hill-climbing Drawbacks

- ❑ Local minima (maxima)
 - ❖ Local, rather than global minima (maxima)
- ❑ Plateau
 - ❖ Area of state space where the evaluation function is essentially flat
 - ❖ The search will conduct a random walk
- ❑ Ridges
 - ❖ Causes problems when states along the ridge are not directly connected - the only choice at each point on the ridge requires uphill (downhill) movement



Diagonal Ridges

4/11/2024

pra-sâmi

25

Problems with Iterative Improvement

□ Given a set pick items to reach a total:

❖ $A = \{15, 5, 8, 2, 12\}$

❖ $Q = 29$

❖ Start $Z = \{15, 5, 8\} = 28$

➢ No further swapping/ additions possible ❌

❖ However, best solution is $\{15, 2, 12\}$

➢ Let's start from right side $Z = \{8, 2, 12\} = 20$

➢ Swap 2 with 15 ; $Z = \{8, 15, 12\} = 35 > 29$

➢ Swap 8 with 15; $Z = \{15, 2, 12\} = 29$


□ Hill Climbing

❖ Start with any maximal subset Z of A

❖ Consider a pair (a_j, a_k) such that $a_j \in Z$ and $a_k \notin Z$.

❖ Replace a_j with a_k in Z in subset sum increases but does not exceed Q .

□ For gradient descent change **increase to decrease**




Gradient Descent

4/11/2024

pra-sâmi

26

Can get stuck in local minima too...



4/11/2024

pra-sâmi

27

Variants of Hill Climbing

- ❑ Escaping shoulder or local optima
 - ❖ May be start with breadth-first search and once we find better optimization function
 - ❖ Start climbing the hill again
- ❑ Prolonged period of exhaustive search
 - ❖ Small period of hill climbing
- ❑ Some what a middle ground between local and systematic search

4/11/2024

pra-sâmi

28

Variants of Hill Climbing

- ❑ What happens if you run into a plateau or ridge
- ❑ Simple solution could be keep moving even if the objective function is same
 - ❖ Impose some limit of moves, if it does not start improving, end the search
 - ❖ But in larger space, we may be shuttling between couple of points
- ❑ Tabu Search
 - ❖ Prevent returning to same state again
 - ❖ Maintain list of states already visited
 - ❖ Fixed length
 - ❖ Add latest visited node, drop the oldest one
 - ❖ In general a list of size 100 to 500 states is sufficient

4/11/2024

pra-sâmi

29

Variants of Hill Climbing

- ❑ Stochastic hill climbing:
 - ❖ Chooses at random from among uphill moves
 - ❖ Converges more slowly, but finds better solutions in some landscapes
- ❑ Different Variations
 - ❖ For each restart: run till end vs. run for fixed time
 - ❖ Run fixed number of restarts vs. run indefinitely
- ❑ Stochastic hill climbing with random walk
 - ❖ Use a probability p
 - ❖ p times use best neighbor
 - ❖ $(1-p)$ time move to a random neighbor
- ❑ As time progresses
 - ❖ Increase p ; more greedy less stochastic


4/11/2024

pra-sâmi

30

Variants of Hill Climbing

- ❑ First-choice hill climbing:
 - ❖ Generate successors randomly until one is better than the current
 - ❖ Good when a state has many successors



4/11/2024



pra-sâmi

31

Variants of Hill Climbing

❑ Random-restart hill climbing:

- ❖ Conducts a series of hill climbing searches from randomly generated initial states, stops when a goal is found
- ❖ It's complete with probability approaching 1
- ❖ Assume each hill climbing search has a probability p of success, then the expected number of restarts required is $1/p$
- ❖ For 8-queen problem, $p = 14\%$, so we need roughly 7 iterations to find a goal
- ❖ Expected # of steps = $\text{cost_to_success} + (1-p)/p * \text{cost_to_failure}$
- ❖ Random-restart hill climbing is very effective for n-queen problem
 - 3 million queens can be solved < 1 min



4/11/2024

pra-sâmi

32

Simulated Annealing

❑ What is Simulated Annealing?

- ❖ *The process used to harden metals and glass by heating them to a high temperature and then gradually cooling them, thus allowing the material to reach a low-energy crystalline state*

❑ Idea

- ❖ Escape local maxima by allowing some "bad" moves but gradually decrease their frequency
- ❖ *'Shake out of the pit'*

❑ One can prove: If T decreases slowly enough, then simulated annealing search will find a global optimum with probability approaching 1

- ❖ Proposed in 1983 by IBM (Kirkpatrick et al) for solving VLSI layout problem
- ❖ Effective in Traveling salesperson, Graph Partitioning, Airline Scheduling, Facility Layout, etc.
- ❖ Later used in Image processing as well.

4/11/2024

pra-sâmi

33

Simulated Annealing

- ❑ Temperature
 - ❖ Let T denote the temperature
 - ❖ Initially Set T to a very high value
 - ❖ Reduce T to Zero gradually with some cooling scheme
 - Reduce it to half after n iterations (exponential)
 - Reduce by x % after n iteration (linear)
- ❑ Algorithm
 - ❖ Initialize T
 - ❖ Set next = randomly selected successor of current
 - Calculate $\Delta E = \text{value (current)} - \text{value (next)}$
 - If $\Delta E > 0$ then set current = next (minimization)
 - Else set current = next with probability $e^{\Delta E/T}$
 - Update T as per the schedule and loop

Boltzmann Function

- ❑ T being in the division, more exploration is permitted at high temperatures
- ❑ More exploitation at , lower temperatures

4/11/2024

pra-sâmi

34

Simulated Annealing : Temperature T

- ❑ High T : probability of **locally bad** move is **high**
- ❑ Low T : probability of **locally bad** move is **low**
- ❑ Depreciate T as time progresses
 - ❖ Over subsequent epochs
 - ❖ Code in a Temperature schedule
 - Reduce 10 % after every 10 epochs

4/11/2024

pra-sâmi

35

Local Beam Search

- ❑ Keeping one node (current state) in the memory!
 - ❖ Is it enough??
 - ❖ Is it not a bit of over reaction
- ❑ Idea:
 - ❖ Keep track of k states rather than just one
 - ❖ Start with k randomly generated states
 - ❖ At each iteration, all the successors of all k states are generated
 - ❖ If anyone is a goal state, stop;
 - ❖ Else select the k best successors from the complete list and repeat

4/11/2024

pra-sâmi

36

Local Beam Search

- ❑ Is it the same as running k random-restart searches in parallel?
 - ❖ No
 - ❖ Useful information is passed among the k parallel search threads
- ❑ Challenge: frequently, all successor end up on same hill
- ❑ Stochastic beam search:
 - ❖ Choose K successors randomly, biased towards good ones
 - ❖ Similar to natural selection, offspring of an organism populate the next generation according to its fitness

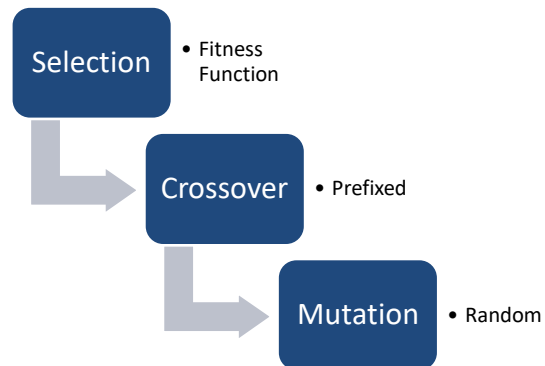
4/11/2024

pra-sâmi

37

Genetic Algorithms

- ❑ Start with k randomly generated states (population)
- ❑ A successor state is generated by combining two parent states
- ❑ Give some representation to state for ease of manipulation
 - ❖ A state is represented as a string over a finite alphabet (often a string of 0s and 1s or digits)
- ❑ Evaluation function (fitness function)
 - ❖ Higher values for better states
- ❑ Produce the next generation of states by selection, crossover, and mutation

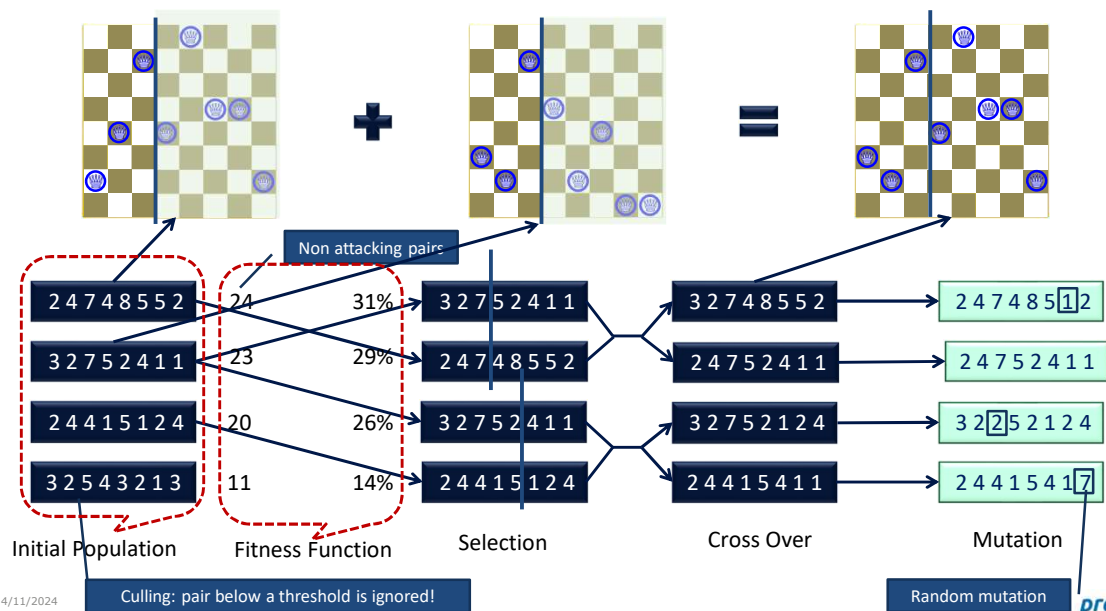


4/11/2024

pra-sâmi

38

Example




4/11/2024

pra-sâmi

39

Fitness Function




- ❑ Fitness function:
 - ❖ Number of non-attacking pairs of queens (min = 0, max = $8 \times 7/2 = 28$)
 - ❖ $24 / (24+23+20+11) = 31\%$
 - ❖ $23 / (24+23+20+11) = 29\%$ etc...
- ❑ Genetic algorithms combine an uphill tendency with random exploration and exchange of information among parallel search threads
- ❑ Advantages come from “crossover”, which raise the level of granularity

4/11/2024

pra-sâmi

40

Gradient Descent



4/11/2024


pra-sâmi

41

Gradient Descent

- ❑ Solve it as a optimization problem
- ❑ Gradient is actually Error Gradient
- ❑ Create some stable loss function
 - ❖ Loss function : $\ell(a, y)$, a in turn is function of W and b
- ❑ Compute gradient $\frac{\partial J}{\partial W}$
- ❑ Update weights $W = W - \alpha \cdot \frac{\partial J}{\partial W}$
- ❑ Similarly $b = b - \alpha \cdot \frac{\partial J}{\partial b}$

Where α is defined as learning rate.




4/11/2024

pra-sâmi

42

Loss / Cost Optimization



4/11/2024

43

Learning Rate : Tough Terrain

- ❑ Finding most optimal gradient descent can be difficult

Question: How to select right learning rate?

- ❑ Too fast and you can miss global minima!
- ❑ Too slow, you can be struck at local minima!
- ❑ Need to look for learning rate that converges smoothly and avoids local minima!
- ❑ Need a learning that 'adapts' to the terrain
- ❑ No Fixed learning rate
- ❑ To change as per the change in gradient
- ❑ Popular algorithms
 - ❖ SGD
 - ❖ Adam
 - ❖ Adadelta
 - ❖ Adagrad
 - ❖ RMSProp

4/11/2024

pra-sâmi

44

Measuring problem-solving performance

The evaluation of a search strategy

- ❑ Completeness:
 - ❖ Is the strategy guaranteed to find a solution when there is one?
- ❑ Optimality:
 - ❖ Does the strategy find the highest-quality solution when there are several different solutions?
- ❑ Time complexity:
 - ❖ How long does it take to find a solution?
- ❑ Space complexity:
 - ❖ How much memory is needed to perform the search?

4/11/2024

pra-sâmi

45

Measuring problem-solving performance

- In AI, complexity is expressed in
 - ❖ b , branching factor, maximum number of successors of any node
 - ❖ d , the depth of the shallowest goal node (depth of the least-cost solution)
 - ❖ m , the maximum length of any path in the state space
- Time and Space is measured in
 - ❖ Number of nodes generated during the search
 - ❖ Maximum number of nodes stored in memory

4/11/2024

pra-sâmi

46

Measuring problem-solving performance

- For effectiveness of a search algorithm
 - ❖ we can just consider the total cost
 - ❖ The total cost = path cost (g) of the solution found + search cost
 - search cost = time necessary to find the solution
- Tradeoff:
 - ❖ *< long time, optimal solution with least g > vs. < shorter time, solution with slightly larger path cost g >*

4/11/2024

pra-sâmi

47

Which method?

- ❑ Exhaustive search for small finite spaces when it is essential that the optimal solution is found
- ❑ A* for medium-sized spaces if heuristic knowledge is available
- ❑ Random search for large evenly distributed homogeneous spaces
- ❑ Hill climbing for discrete spaces where a sub-optimal solution is acceptable


4/11/2024

pra-sâmi

48

Parallel Depth First Search:

- ❑ The critical issue in parallel depth-first search algorithms is the distribution of the search space among the processors



4/11/2024

pra-sâmi

49

Reflect...

- ❑ In heuristic hill climbing , paths typically not retained - very little memory needed
- ❑ Depending on initial state, can get stuck in local maxima/minima
 - ❖ Right choice will take you to global maxima
 - ❖ Making some “bad” choices is actually not that bad
- ❑ Variants of Hill Climbing
 - ❖ Stochastic
 - ❖ First-choice
 - ❖ Random-restart
- ❑ In many optimization problems, the path to the goal is irrelevant; the goal state itself is the solution

4/11/2024

pra-sâmi

50

Reflect

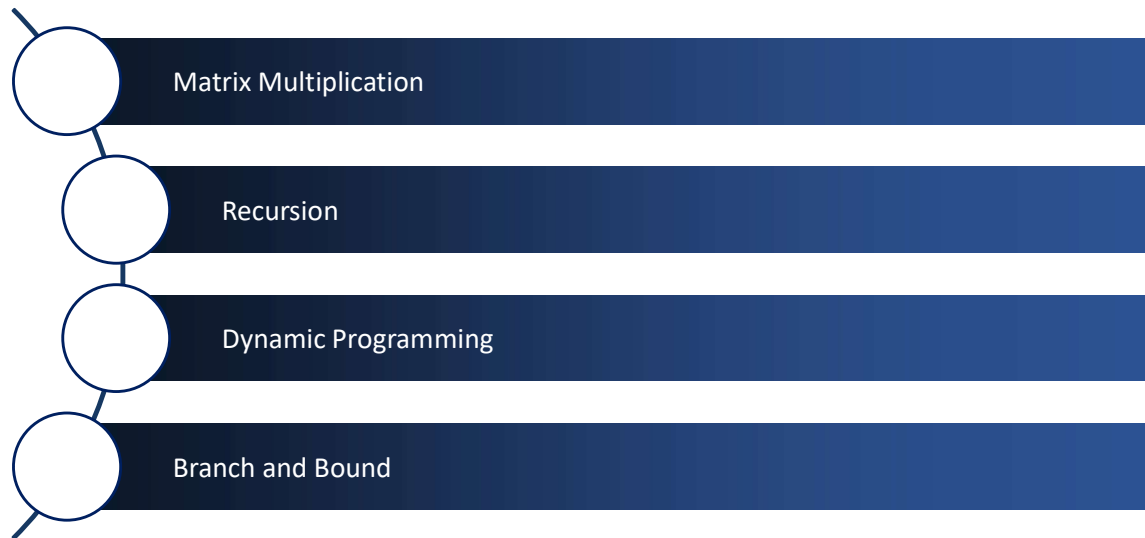
- ❑ We are looking at AI problems as a search problems
 - ❖ Defined the problem as State Space Search
 - ❖ Looked at solutions as:
 - Path to Goal state
 - Or Goal as some better state
- ❑ We looked at:
 - ❖ Informed search strategies
 - Greedy best-first search
 - A* search
 - Memory Bound Search
 - ❖ Blind Search
 - Breadth-first search
 - Depth-first search
 - Bidirectional search
 - ❖ Local search strategies- Hill climbing, simulated annealing.
 - ❖ Performance measurement

4/11/2024

pra-sâmi

51

Next Session



4/11/2024

pra-sâmi

52



4/11/2024

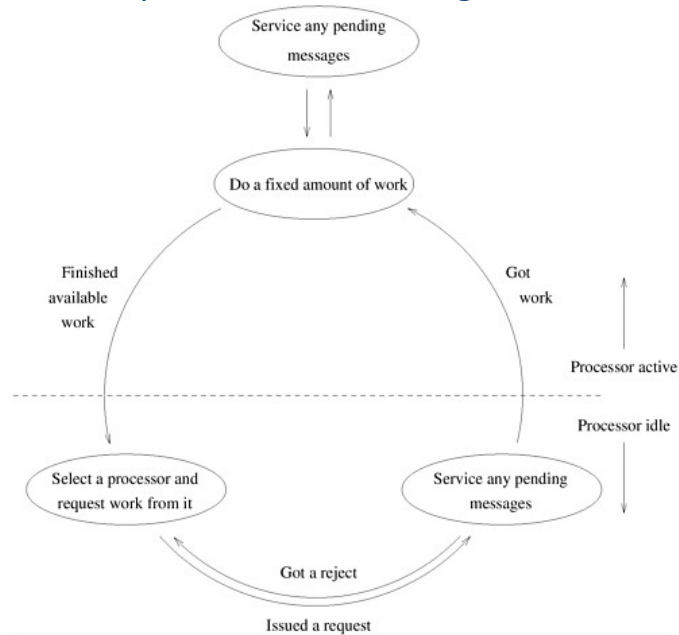
pra-sâmi

ADDITIONAL MATERIAL

pra-sâmi

55

A generic scheme for dynamic load balancing



4/11/2024

pra-sâmi

56

Parallel Depth First Search:

- ❑ Important Parameters of Parallel DFS
- ❑ Load -Balancing Schemas:
 - ❖ Asynchronous Round Robin
 - ❖ Global Round Robin
 - ❖ Random Polling

4/11/2024

pra-sâmi

57

Parallel Best first Search:

- ❑ This algorithm contains two main components:
 - ❖ Open list , Close list
- ❑ In most parallel formulations of BFS, different processors concurrently expand different nodes from the open list
- ❑ There are two problems with this approach:
 - ❖ The termination criterion of sequential BFS fails for parallel BFS
 - ❖ Since the open list is accessed for each node expansion, it must be easily accessible to all processors, which can severely limit performance


4/11/2024

pra-sâmi

58

Parallel Best first Search:

- A general schematic for parallel best-first search using a centralized strategy



59

Parallel Best first Search:

- One way to avoid the contention due to a centralized open list is to let each processor have a local open list.
- The processors must communicate among themselves to minimize unnecessary search

4/11/2024

pra-sâmi

60

Parallel Best first Search:

- ❑ Communication Strategies for Parallel Best-First Tree Search
 - ❖ random communication strategy
 - ❖ ring communication strategy
 - ❖ blackboard communication strategy


4/11/2024

pra-sâmi

61

Parallel Binary Search Algorithm:

- ❑ We have an ordered array ,we have two processors
- ❑ We part our array to $P+1$ parts , where p is number of processors



4/11/2024

pra-sâmi