



# AGENTS

Fundamentals of Artificial Intelligence

Session 13

Pramod Sharma  
pramod.sharma@prasami.com

2

## Agenda




4/14/2024

*pra-sâmi*

3


## Mars Rover Perseverance



Source: Nasa

BBC

4/14/2024



pra-sâmi

4

## Acting Rationally

	Human-like Intelligence	"Ideal" Intelligent/ Pure Rationality
Thought/ Reasoning ("modeling thought / brain)	2. Thinking humanly	3. Thinking Rationally
Behavior/Actions "behaviorism" "mimics behavior"	1. Acting Humanly	4. Acting Rationally

- ❑ In the "laws of thought" approach to AI, the emphasis was on correct inferences
  - ❖ Making correct inferences is sometimes part of being a rational agent
  - ❖ But there are reflex action....
- ❑ The rational-agent approach has advantages
  - ❖ More general than the "laws of thought" approach
  - ❖ The standard of rationality is mathematically well defined
  - ❖ But computational demands are just too high

4/14/2024

pra-sâmi

5

## Agent

- ❑ We sense the world around us
  - ❖ Create notion of a world in our head
  - ❖ Convert it to symbols
- ❑ These symbols and their interaction we try to define by known law of physics
- ❑ That's the idea of Agent
- ❑ Agent: anything that can be viewed as:
  - ❖ **Perceiving** its environment through **sensors** and
  - ❖ **Acting** upon that **environment** through **actuators**

4/14/2024

pra-sâmi

6

## Agent

- ❑ Human agent has:
  - ❖ eyes, ears, and other organs for sensors and hands, legs, vocal tract, and so on for actuators
- ❑ Robotic agent might have:
  - ❖ cameras and infrared range finders for sensors and various motors for actuators
- ❑ Software agent receives:
  - ❖ keystrokes, file contents, and network packets as sensory inputs and acts on the environment by displaying on the screen, writing files, and sending network packets
- ❑ Percept: The agent's perceptual inputs at any given instant
- ❑ Percept Sequence : complete history of everything the agent has ever perceived.
  - ❖ An agent's choice of action at any given instant can depend on the entire percept sequence observed to date, **but not on anything it hasn't perceived**
- ❑ Agent Function: describes agent's behavior that maps any given percept sequence to an action

4/14/2024

pra-sâmi

12

## Rational Agent

- ❑ Has clear preference
  - ❖ Action for every possible sequence is clearly programmed
  - ❖ Every entry in the table for the agent function is filled out correctly
- ❑ Acts in a way to maximize its performance measure with all possible actions
- ❑ A rational agent is said to perform the right things
- ❑ AI is about creating rational agents to use for various real-world scenarios
- ❑ For an AI agent, the rational action is most important
  - ❖ For each best possible action, agent gets the positive reward
  - ❖ For each wrong action, an agent gets a negative reward

4/14/2024

pra-sâmi

13

## Rational Agent Models Uncertainty

- ❑ A rational agent not only gathers information but also learns as much as possible from what it perceives
- ❑ The agent's initial configuration could reflect some prior knowledge of the environment, but as the agent gains experience this may be modified and augmented.
- ❑ An agent that does not perceive or learn is fragile!
- ❑ Dung Beetle
  - ❖ A lowly dung beetle, after digging its nest and laying its eggs, it fetches a ball of dung from a nearby heap to plug the entrance.
  - ❖ If the ball of dung is removed from its grasp en route, the beetle continues its task and pantomimes plugging the nest with the nonexistent dung ball, never noticing that it is missing
- ❑ SpheX Wasp ( slightly more intelligent)
  - ❖ The female sphex will dig a burrow, go out and sting a caterpillar and drag it to the burrow, enter the burrow again to check all is well, drag the caterpillar inside, and lay its eggs.
  - ❖ The caterpillar serves as a food source when the eggs hatch.
  - ❖ So far so good, but if an entomologist moves the caterpillar a few inches away while the sphex is doing the check, it will revert to the "drag" step of its plan and will continue the plan without modification, even after dozens of caterpillar-moving interventions.
  - ❖ The sphex is unable to learn that its innate plan is failing, and thus will not change it.

4/14/2024

pra-sâmi

15

## Distinguish between Rationality and Omniscience

- ❑ Rational behavior - maximizes expected performance
  - ❖ Look both ways while crossing the road, cross if no approaching vehicle
- ❑ Omniscience (state of knowing everything) – maximizes actual performance
  - ❖ Person crossing a road killed by flying hoarding? (Rational but not Omniscience)
  - ❖ Impossible to design agent for Omniscience
- ❑ Definition of rationality does not require omniscience
  - ❖ Because the rational choice depends only on the percept sequence to date
- ❑ Doing actions in order to modify future percepts—sometimes called information gathering—is an important part of rationality (**Exploration**)

4/14/2024

pra-sâmi

16

## Rational Agents

- ❑ An autonomous entity :
  - ❖ Act upon an environment using sensors and actuators for achieving goals
  - ❖ May learn from the environment to achieve their goals
  - ❖ A thermostat is an example
- ❑ Following are the main four rules for an AI agent:
  - ❖ Rule 1: An AI agent must have the ability to perceive the environment
  - ❖ Rule 2: The observation must be used to make decisions
  - ❖ Rule 3: Decision should result in an action
  - ❖ Rule 4: The action taken by an AI agent must be a rational action
- ❑ Well-behaved agents
  - ❖ For **each possible percept sequence**, a rational agent should select an action that is **expected to maximize its performance measure**, given **the evidence provided** by the percept sequence and **whatever built-in knowledge** the agent has

4/14/2024

pra-sâmi

17

## Type of agents

- ❑ Five classes based on their degree of perceived intelligence and capability
- ❑ All these agents can improve their performance and generate better action over the time
  - ❖ Simple Reflex Agent
  - ❖ Model-based Reflex Agent
  - ❖ Goal-based Agents
  - ❖ Utility-based Agent
  - ❖ Learning Agent

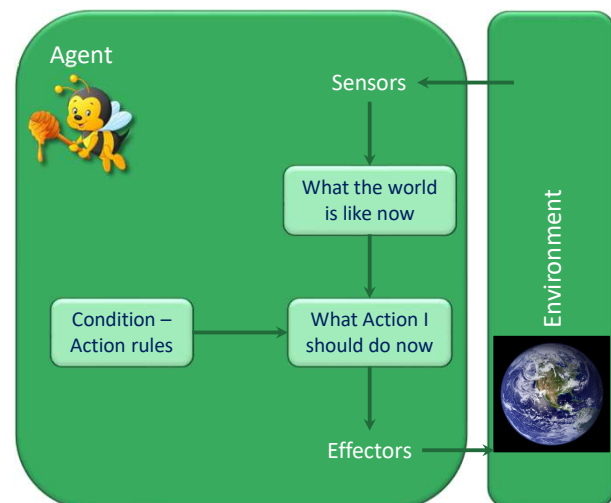
4/14/2024

pra-sâmi

18

## Simple Reflex Agent

- ❑ Agents take decisions on the basis of the current percepts and ignore the rest of the percept history
- ❑ Successful in the fully observable environment only
- ❑ No percept history considered
- ❑ Works on “Condition – Action” rule
  - ❖ Maps the current state to action
    - If hand is in fire, pull it out
    - Reflex action / muscle memory similar to humans
  - ❖ Example: Room Cleaner agent
    - There is dirt in the room → pick it up.


<https://www.doc.ic.ac.uk/project/examples/2005/163/g0516334/sra.html>

4/14/2024

pra-sâmi

19

## Limitation of Simple Reflex agent

- ❑ They have very limited intelligence
- ❑ They do not have knowledge of non-perceptual parts of the current state
  - ❖ Take example of a mars lander designed to collect a rocks
  - ❖ A simple reflex agent will keep collecting rocks even if same rock is found again
  - ❖ It doesn't take into account that it already picked up this type of rock
- ❑ Not adaptive to changes in the environment

4/14/2024

pra-sâmi

20

## Model-based Reflex Agent

- ❑ Can work in a partially observable environment, and track the situation
- ❑ Two important factors:
  - ❖ Model: It is knowledge about "how things happen in the world," so it is called a Model-based agent
  - ❖ Internal State: It is a representation of the current state based on percept history
- ❑ These agents have the model, "which is knowledge of the world" and based on the model they perform actions
  - ❖ Mars Lander after picking up its first sample, it stores this in the internal state of the world around it so when it come across the second same sample it passes it by and saves space for other samples
- ❑ Updating the agent state requires information about:
  - ❖ How the world evolves
  - ❖ How the agent's action affects the world
    - If our mars Lander took a sample under a precarious ledge it could displace a rock and it could be crushed

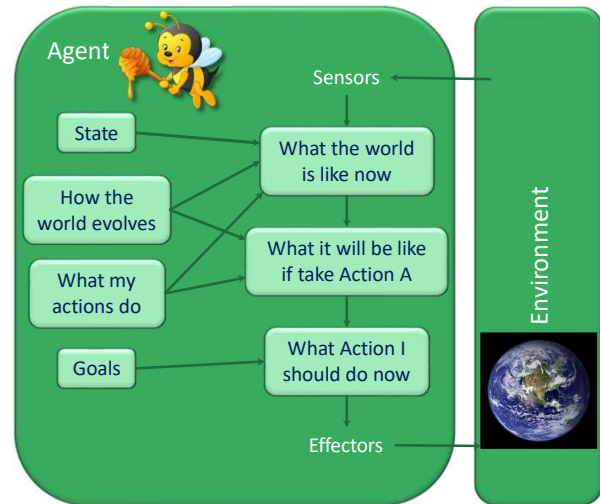
4/14/2024

pra-sâmi

21

## Goal Based Agents

- ❑ Agent having a set of goals with desirable situations
- ❑ Agent uses these goals with a set of actions to move towards the goal
  - ❖ Predicted outcomes to see which action(s) achieve our goal(s) better
  - ❖ For achieving the goals, can take an action or many actions
- ❑ Two subfields devoted to finding sequences of actions to achieve goals
  - ❖ Search
  - ❖ Planning
- ❑ Before acting the agent:
  - ❖ Reviews many actions
  - ❖ Chooses the one which come closest to achieving its goals
  - ❖ Whereas the reflex agents just have an automated response for the situations




4/14/2024

pra-sâmi

22

## Goal based agents are Always learning

- ❑ Much more flexible because the knowledge used for decision making is represented explicitly and can be modified
- ❑ For example if our mars Lander needed to get up a hill the agent can update its knowledge on how much power to put into the wheels to gain certain speeds, through this all relevant behaviors will now automatically follow the new knowledge on moving
- ❑ Still these agents are not measuring and storing their efficiency



- ❑ Please welcome Utility Based agents

4/14/2024

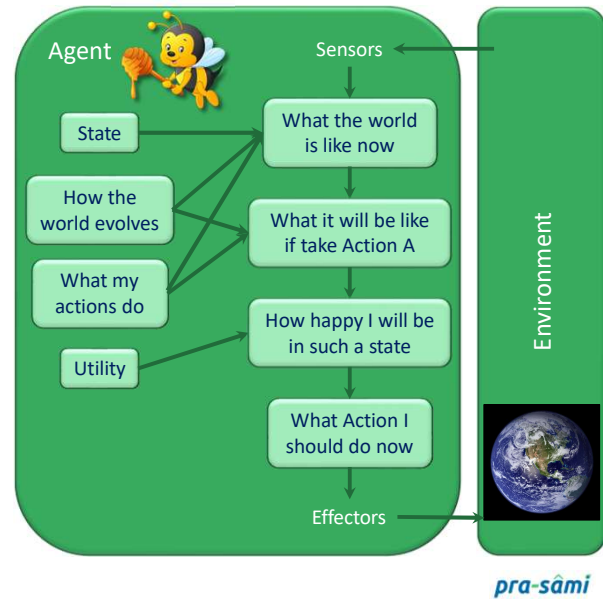
pra-sâmi



23

## Utility-based Agents

- ❑ Just having goals is not enough, the agents need to also achieve these goals in best possible way
- ❑ They an extra component of utility measurement
  - ❖ Utilization provides a measure of success at a given state
  - ❖ E.g. how many steps to achieve the goal
- ❑ Utility-based agent act based not only goals but also the best way to achieve the goal
- ❑ The Utility-based agent is useful when there are multiple possible alternatives, and an agent has to choose in order to perform the best action




4/14/2024

24

## Utility-based Agents

- ❑ The utility function maps each state to a real number to check how efficiently each action achieves the goals
- ❑ For example let's show our mars Lander on the surface of mars with an obstacle in its way.
  - ❖ In a goal based agent it is uncertain which path will be taken by the agent
  - ❖ some are clearly not as efficient as others
  - ❖ Utility based agent the best path will have the best output from the utility function and that path will be chosen

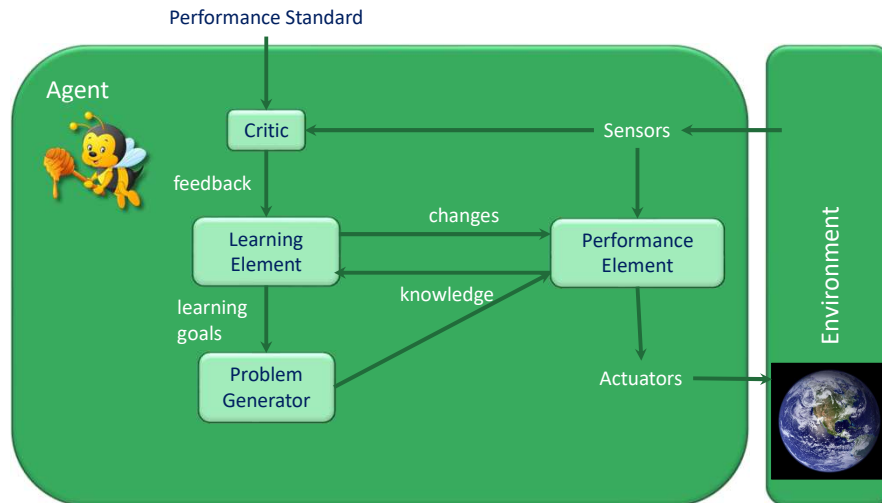


4/14/2024

25

## A Learning Agent

- Has four broad parts



4/14/2024

pra-sâmi

29

## Structure of an AI Agent

- Architecture: Architecture is machinery that an AI agent executes on
- Agent Function: Agent function is used to map a percept to an action
- Agent program: Agent program is an implementation of agent function. An agent program executes on the physical architecture to produce function  $f$
- PEAS Representation
  - ❖ **P**: Performance measure
  - ❖ **E**: Environment
  - ❖ **A**: Actuators
  - ❖ **S**: Sensors

4/14/2024

pra-sâmi

30

## PEAS

- ❑ Performance:
  - ❖ Safety, time Legal drive and comfort
- ❑ Environment:
  - ❖ Roads, other cars, pedestrians, road signs and everything else
- ❑ Actuators:
  - ❖ Steering, accelerator, brake, signal, horn
- ❑ Sensors:
  - ❖ Camera, sonar, GPS, Speedometer, odometer, accelerometer, engine sensors, keyboard



4/14/2024

pra-sâmi

31

## PEAS

- ❑ Performance:
  - ❖ Cleanness, efficiency: distance traveled to clean, battery life, security
- ❑ Environment:
  - ❖ Room, table, wood floor, carpet, different obstacles
- ❑ Actuators:
  - ❖ Wheels, different brushes, vacuum extractor
- ❑ Sensors:
  - ❖ Camera, dirt detection sensor, cliff sensor, bump sensors, infrared wall sensors.



iRobot Roomba 675

4/14/2024

pra-sâmi

32

## Reflect

#	Agent	Example
1	Simple Reflex Agent	<ul style="list-style-type: none"> <li>Thermostat,</li> <li>Atari console with left/ right move and fire button</li> </ul>
2	Model-based Reflex Agent	<ul style="list-style-type: none"> <li>Traffic Camera taking pic of speeding vehicle</li> <li>Atari console forcing you to remain within a lane while you drive</li> </ul>
3	Goal Based Agents	<ul style="list-style-type: none"> <li>Google's Waymo driverless cars</li> <li>Programmed with an end destination, or goal</li> <li>The car will then "think" and make the right decisions in order to reach the destination</li> <li>Still only thinking of Goal. Not too worried about time it is taking to reach there</li> </ul>
4	Utility-based Agents	<ul style="list-style-type: none"> <li>Route recommendation system which solves the 'best' route to reach a destination</li> </ul>
5	A Learning Agent	<ul style="list-style-type: none"> <li>Human : Can learn to ride a bicycle, even though, at birth, no human possesses this skill</li> </ul>


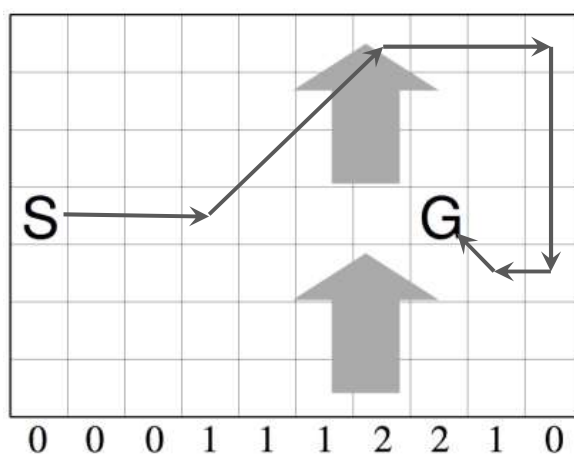
4/14/2024

pra-sâmi

33

## Learning Agent: Windy Gridworld Example

- Agent can move one step at a time
- Reward = -1 per time-step until reaching goal



4/14/2024

pra-sâmi

34

## Properties of Task Environments

- ❑ Fully observable vs. partially observable
  - ❖ How much can we observe
- ❑ Single agent vs. multi-agent
- ❑ Deterministic vs. stochastic
- ❑ Episodic vs. sequential
- ❑ Static vs. dynamic
- ❑ Discrete vs. continuous
- ❑ Known vs. unknown
  - ❖ Is physics around the environment known? (Chandrayaan?)

4/14/2024

pra-sâmi

35

## Rented car – new country – unfamiliar geography and traffic laws

Task Environment	Observable	Agents	Deterministic	Episodic	Static	Discrete
Crossword puzzle	Fully	Single	Deterministic	Sequential	Static	Discrete
Chess with a clock	Fully	Multi	Deterministic	Sequential	Semi	Discrete
Poker	Partially	Multi	Stochastic	Sequential	Static	Discrete
Backgammon	Fully	Multi	Stochastic	Sequential	Static	Discrete
Taxi driving	Partially	Multi	Stochastic	Sequential	Dynamic	Continuous
Medical diagnosis	Partially	Single	Stochastic	Sequential	Dynamic	Continuous
Image analysis	Fully	Single	Deterministic	Episodic	Semi	Continuous
Part-picking robot	Partially	Single	Stochastic	Episodic	Dynamic	Continuous
Refinery controller	Partially	Single	Stochastic	Sequential	Dynamic	Continuous
Interactive English tutor	Partially	Multi	Stochastic	Sequential	Dynamic	Discrete

4/14/2024

pra-sâmi

36

## Reflect...

- ❑ An agent is something that perceives and acts in an environment.
- ❑ The agent function for an agent specifies the action taken by the agent in response to any percept sequence
- ❑ The performance measure evaluates the behavior of the agent in an environment.
- ❑ A rational agent acts so as to maximize the expected value of the performance measure, given the percept sequence it has seen so far
- ❑ A task environment specification includes the performance measure, the external environment, the actuators, and the sensors.
- ❑ In designing an agent, the first step must always be to specify the task environment as fully as possible

4/14/2024

pra-sâmi

37

## Reflect...

- ❑ Task environments vary along several significant dimensions.
- ❑ Environments can be fully or partially observable, single-agent or multi-agent, deterministic or stochastic, episodic or sequential, static or dynamic, discrete or continuous, and known or unknown
- ❑ The agent program implements the agent function. There exists a variety of basic agent-program designs reflecting the kind of information made explicit and used in the decision process.
- ❑ The designs vary in efficiency, compactness, and flexibility. The appropriate design of the agent program depends on the nature of the environment
- ❑ Simple reflex agents respond directly to percepts, whereas model-based reflex agents maintain internal state to track aspects of the world that are not evident in the current percept.
- ❑ Goal-based agents act to achieve their goals, and utility-based agents try to maximize their own expected "happiness."
- ❑ All agents can improve their performance through learning

4/14/2024

pra-sâmi

38

## Reflect...

- ❑ \_\_\_\_\_ are used for perceiving and \_\_\_\_\_ are used for acting upon the environment?
  - ❖ Sensors and Actuators
  - ❖ Sensors
  - ❖ Perceiver
  - ❖ None of the above
- ❑ What is meant by agent's percept sequence?
  - ❖ Used to perceive the environment
  - ❖ Complete history of actuator
  - ❖ Complete history of perceived things
  - ❖ None of the above
- ❑ What is the function of an artificial intelligence "Agent"?
  - ❖ Mapping of goal sequence to an action
  - ❖ Work without the direct interference of the people
  - ❖ Mapping of precept sequence to an action
  - ❖ Mapping of environment sequence to an action
- ❑ What is the rule of simple reflex agent?
  - ❖ Simple-action rule
  - ❖ Condition-action rule
  - ❖ Simple & Condition-action rule
  - ❖ None of the above

4/14/2024

pra-sâmi

39

## Reflect...

- ❑ The composition for agents in artificial intelligence are \_\_\_\_\_
  - ❖ Program only
  - ❖ Architecture only
  - ❖ Both Program and Architecture
  - ❖ None of the above
- ❑ In which agent does the problem generator is present?
  - ❖ Learning agent
  - ❖ Observing agent
  - ❖ Reflex agent
  - ❖ None of the above
- ❑ Which agent deals with happy and unhappy states?
  - ❖ Simple reflex agent
  - ❖ Model based agent
  - ❖ Learning agent
  - ❖ Utility based agent
- ❑ What is an 'agent'?
  - ❖ Perceives its environment through sensors and acting upon that environment through actuators
  - ❖ Takes input from the surroundings and uses its intelligence and performs the desired operations
  - ❖ A embedded program controlling line following robot
  - ❖ All of the above

4/14/2024

pra-sâmi

40

## Reflect...

- ❑ Agents behavior can be best described by \_\_\_\_\_
  - ❖ Perception sequence
  - ❖ Agent function
  - ❖ Sensors and Actuators
  - ❖ Environment in which agent is performing
- ❑ Rational agent is the one who always does the right thing.
  - ❖ True
  - ❖ False
- ❑ Performance Measures are fixed for all agents.
  - ❖ True
  - ❖ False
- ❑ An omniscient agent knows the actual outcome of its actions and can act accordingly; but omniscience is impossible in reality
  - ❖ True
  - ❖ False
- ❑ Rational Agent always does the right thing; but Rationality is possible in reality.
  - ❖ True
  - ❖ False
- ❑ The Task Environment of an agent consists of \_\_\_\_\_
  - ❖ Sensors
  - ❖ Actuators
  - ❖ Performance Measures
  - ❖ All of the mentioned

4/14/2024

pra-sâmi

41

## Reflect...

- ❑ Categorize Crossword puzzle in Fully Observable / Partially Observable.
  - ❖ Fully Observable
  - ❖ Partially Observable
  - ❖ All of the mentioned
  - ❖ None of the mentioned
- ❑ The game of Poker is a single agent.
  - ❖ True
  - ❖ False
- ❑ What is called an exploration problem?
  - ❖ State and actions are unknown to the agent
  - ❖ State and actions are known to the agent
  - ❖ Only actions are known to agent
  - ❖ None of the mentioned

4/14/2024

pra-sâmi



43



4/14/2024

*pra-sâmi*

## EXTRA MATERIAL

Details of different environments

*pra-sâmi*

45

## Properties of Task Environments

- ❑ Range of task environments is vast;
- ❑ Need to reduce number of dimensions to some manageable number
- ❑ What dimensions to consider?

4/14/2024

pra-sâmi

46

## Fully Observable vs. Partially Observable

- ❑ Fully observable
  - ❖ An agent's sensors give it access to the complete state of the environment at each point in time
  - ❖ The sensors detect all aspects that are relevant to the choice of action; relevance, in turn, depends on the performance measure
- ❑ Fully observable environments are convenient because the agent need not maintain any internal state to keep track of the world
- ❑ An environment might be partially observable because:
  - ❖ Noisy and inaccurate sensors
  - ❖ Parts of the state are simply missing from the sensor data
- ❑ Examples
  - ❖ A vacuum agent with only a local dirt sensor cannot tell whether there is dirt in other squares
  - ❖ An automated taxi cannot see what other drivers are thinking
- ❑ The agent has no sensors at all then the environment is unobservable
  - ❖ The agent's goals may still be achievable, sometimes with certainty

4/14/2024

pra-sâmi

47

## Single Agent vs. Multi-agent

- ❑ Fairly Straight forward
  - ❖ An agent solving a crossword puzzle by itself is clearly in a single-agent environment, whereas an agent playing chess is in a two agent environment
- ❑ In most cases we struggle to classify
  - ❖ How an entity may be viewed as an agent?
  - ❖ Do we know which entities must be viewed as agents.
  - ❖ Does an agent A (the taxi driver for example) have to treat an object B (another vehicle) as an agent, or can it be treated merely as an object behaving according to the laws of physics, analogous to waves at the beach or leaves blowing in the wind?
  - ❖ Are they part of the environment?
- ❑ The key distinction:
  - ❖ Whether B's behavior is best described as maximizing a performance measure whose value depends on agent A's behavior.

4/14/2024

pra-sâmi

48

## Single Agent vs. Multi-agent

- ❑ Competitive multi-agent environment – chess
  - ❖ The opponent entity B is trying to maximize its performance measure,
  - ❖ Minimizes agent A's performance measure
- ❑ Partially cooperative multi-agent environment - taxi-driving environment
  - ❖ Avoiding collisions maximizes the performance measure of all agents
  - ❖ It is also partially competitive because, for example, only one car can occupy a parking space.
- ❑ The agent-design problems in multi-agent environments are often quite different from those in single-agent environments
  - ❖ Communication often emerges as a rational behavior in multi-agent environments;
  - ❖ In some competitive environments, randomized behavior is rational because it avoids the pitfalls of predictability

4/14/2024

pra-sâmi

49

## Deterministic vs. Stochastic

- ❑ If the next state of the environment is completely determined by:
  - ❖ The current state and the action executed by the agent → the environment is deterministic; otherwise, it is stochastic
  - ❖ The Toy Vacuum World is deterministic
- ❑ Most real situations are so complex that it is impossible to keep track of all the unobserved aspects;
  - ❖ Must be treated as stochastic.
  - ❖ Taxi driving is clearly stochastic in this sense, because one can never predict the behavior of traffic exactly; moreover, one's tires blow out and one's engine seizes up without warning.
  - ❖ Variations can include stochastic elements such as randomly appearing dirt and an unreliable suction mechanism
- ❑ An environment is uncertain if it is not fully observable or not deterministic
- ❑ In our case, a non-deterministic environment is one in which actions are characterized by their possible outcomes, but no probabilities are attached to them
  - ❖ Behavior of other drives are non-deterministic.
- ❑ Nondeterministic environment descriptions are usually associated with performance measures that require the agent to succeed for all possible outcomes of its actions.

4/14/2024

pra-sâmi

50

## Episodic vs. Sequential

- ❑ Episodic task environment – the agent's experience is divided into atomic episodes
  - ❖ In each episode the agent receives a percept and then performs a single action or sequence of actions
  - ❖ The next episode does not depend on the actions taken in previous episodes
  - ❖ Many classification tasks are episodic
- ❑ For example
  - ❖ An agent that has to spot defective parts on an assembly line bases each decision on the current part, regardless of previous decisions
  - ❖ The current decision doesn't affect whether the next part is defective
- ❑ Sequential environments – the current decision could affect all future decisions
  - ❖ Chess and taxi driving are sequential: in both cases, short-term actions can have long-term consequences
- ❑ Episodic environments are much simpler than sequential environments because the agent does not need to think ahead.

4/14/2024

pra-sâmi

51

## Static vs. Dynamic

- ❑ If the environment can change while an agent is deliberating, then we say the environment is dynamic for that agent; otherwise, it is static
- ❑ Static environments are easy to deal with because the agent need not keep looking at the world while it is deciding on an action, nor need it worry about the passage of time
- ❑ Dynamic environments are continuously asking the agent what it wants to do; if it hasn't decided yet, that counts as deciding to do nothing
- ❑ If the environment itself does not change with the passage of time but the agent's performance score does, then we say the environment is semi-dynamic
- ❑ Taxi driving is clearly dynamic: the other cars and the taxi itself keep moving while the driving algorithm dithers about what to do next
- ❑ Chess, when played with a clock, is semi-dynamic
- ❑ Crossword puzzles are static

4/14/2024

pra-sâmi

52

## Discrete vs. Continuous

- ❑ The discrete/continuous distinction applies to the state of the environment
  - ❖ The way time is handled,
  - ❖ To the percepts and actions of the agent
- ❑ For example, the chess environment has a finite number of distinct states
  - ❖ Chess also has a discrete set of percepts and actions
- ❑ Taxi driving is a continuous-state and continuous-time problem:
  - ❖ the speed and location of the taxi and of the other vehicles sweep through a range of continuous values and do so smoothly over time
  - ❖ Taxi-driving actions are also continuous (steering angles, etc.)
- ❑ Input from digital cameras is discrete, strictly speaking, but is typically treated as representing continuously varying intensities and locations.

4/14/2024

pra-sâmi

53

## Known vs. Unknown

- ❑ Strictly speaking, this distinction refers not to the environment itself but to the agent's (or designer's) state of knowledge about the "laws of physics" of the environment
- ❑ Known Environment – the outcomes (or outcome probabilities if the environment is stochastic) for all actions are given
- ❑ Unknown Environment – the agent will have to learn how it works in order to make good decisions
- ❑ The distinction between known and unknown environments is not the same as the one between fully and partially observable environments
- ❑ It is quite possible for a known environment to be partially observable—for example, in solitaire card games
  - ❖ Rules are known, still unable to see the cards that have not yet been turned over
  - ❖ Conversely, an unknown environment can be fully observable – in a new video game, the screen may show the entire game state but I still don't know what the buttons do until I try them

4/14/2024

*pra-sâmi*